

```
In [1]: import numpy as np
import pandas as pd
import os
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```


```
In [2]: data = pd.read_csv("breast_cancer_data.csv")
```

```
In [3]: data.head(10)
```

```
Out[3]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothn
0	842302	M	17.99	10.38	122.80	1001.0	
1	842517	M	20.57	17.77	132.90	1326.0	
2	84300903	M	19.69	21.25	130.00	1203.0	
3	84348301	M	11.42	20.38	77.58	386.1	
4	84358402	M	20.29	14.34	135.10	1297.0	
5	843786	M	12.45	15.70	82.57	477.1	
6	844359	M	18.25	19.98	119.60	1040.0	
7	84458202	M	13.71	20.83	90.20	577.9	
8	844981	M	13.00	21.82	87.50	519.8	
9	84501001	M	12.46	24.04	83.97	475.9	

10 rows × 33 columns



```
In [4]: data.drop(['Unnamed: 32', "id"], axis=1, inplace=True)
data.diagnosis = [1 if each == "M" else 0 for each in data.diagnosis]
y = data.diagnosis.values
x_data = data.drop(['diagnosis'], axis=1)
```

```
In [7]: x_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 30 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   radius_mean                           569 non-null    float64
1   texture_mean                           569 non-null    float64
2   perimeter_mean                         569 non-null    float64
3   area_mean                             569 non-null    float64
4   smoothness_mean                        569 non-null    float64
5   compactness_mean                       569 non-null    float64
6   concavity_mean                         569 non-null    float64
7   concave points_mean                    569 non-null    float64
8   symmetry_mean                          569 non-null    float64
9   fractal_dimension_mean                 569 non-null    float64
10  radius_se                              569 non-null    float64
11  texture_se                             569 non-null    float64
12  perimeter_se                           569 non-null    float64
13  area_se                                569 non-null    float64
14  smoothness_se                          569 non-null    float64
15  compactness_se                         569 non-null    float64
16  concavity_se                           569 non-null    float64
17  concave points_se                      569 non-null    float64
18  symmetry_se                            569 non-null    float64
19  fractal_dimension_se                   569 non-null    float64
20  radius_worst                           569 non-null    float64
21  texture_worst                          569 non-null    float64
22  perimeter_worst                        569 non-null    float64
23  area_worst                             569 non-null    float64
24  smoothness_worst                       569 non-null    float64
25  compactness_worst                      569 non-null    float64
26  concavity_worst                        569 non-null    float64
27  concave points_worst                   569 non-null    float64
28  symmetry_worst                         569 non-null    float64
29  fractal_dimension_worst                 569 non-null    float64
dtypes: float64(30)
memory usage: 133.5 KB

```

```
In [8]: from sklearn.preprocessing import MinMaxScaler
```

```
In [9]: scaler = MinMaxScaler(feature_range=(0, 1))
```

```
In [10]: x_scaled=scaler.fit_transform(x_data)
```

```
In [12]: x_scaled
```

```
Out[12]: array([[0.52103744, 0.0226581 , 0.54598853, ..., 0.91202749, 0.59846245,
                0.41886396],
                [0.64314449, 0.27257355, 0.61578329, ..., 0.63917526, 0.23358959,
                0.22287813],
                [0.60149557, 0.3902604 , 0.59574321, ..., 0.83505155, 0.40370589,
                0.21343303],
                ...,
                [0.45525108, 0.62123774, 0.44578813, ..., 0.48728522, 0.12872068,
                0.1519087 ],
                [0.64456434, 0.66351031, 0.66553797, ..., 0.91065292, 0.49714173,
                0.45231536],
                [0.03686876, 0.50152181, 0.02853984, ..., 0.          , 0.25744136,
                0.10068215]])
```


```
In [17]: scaled_data = pd.DataFrame(x_scaled, columns=x_data.columns)
```

```
In [18]: scaled_data
```

```
Out[18]:
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean
0	0.521037	0.022658	0.545989	0.363733	0.593753	0.016299
1	0.643144	0.272574	0.615783	0.501591	0.289880	0.015858
2	0.601496	0.390260	0.595743	0.449417	0.514309	0.015711
3	0.210090	0.360839	0.233501	0.102906	0.811321	0.016708
4	0.629893	0.156578	0.630986	0.489290	0.430351	0.015454
...	...	...	...	...	...	...
564	0.690000	0.428813	0.678668	0.566490	0.526948	0.016009
565	0.622320	0.626987	0.604036	0.474019	0.407782	0.016299
566	0.455251	0.621238	0.445788	0.303118	0.288165	0.015858
567	0.644564	0.663510	0.665538	0.475716	0.588336	0.016708
568	0.036869	0.501522	0.028540	0.015907	0.000000	0.016299

569 rows × 30 columns



## splitting train and test samples

```
In [19]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(scaled_data, y, test_size=0.15,
```

```
In [20]: print("x train: ",x_train.shape)
print("x test: ",x_test.shape)
print("y train: ",y_train.shape)
print("y test: ",y_test.shape)
```

```

x train: (483, 30)
x test: (86, 30)
y train: (483,)
y test: (86,)

```

## logistic model interpretation

```

In [23]: def sigmoid(z):
          y_head = 1/(1+np.exp(-z))
          return y_head

```

```

In [28]: def compute_loss(y_true, y_pred):
          # Avoiding Log(0) with small addition
          val = 1e-15
          y_pred = np.clip(y_pred, val, 1 - val)
          return -np.mean(y_true * np.log(y_pred) + (1 - y_true) * np.log(1 - y_pred))

```

```

In [47]: # Gradient descent optimization for Logistic regression
def logistic_regression(X, y, learning_rate=0.01, epochs=1000):
    """
    X: Input feature matrix (shape: m x n)
    y: Target labels (shape: m x 1)
    learning_rate: Step size for gradient descent
    epochs: Number of iterations
    """
    # Initialize weights and bias
    m, n = X.shape # m: number of samples, n: number of features
    weights = np.zeros(n)
    bias = 0
    loss_list = []
    loss_list2 = []
    index = []
    for epoch in range(epochs):
        # Compute linear combination
        z = np.dot(X, weights) + bias

        # Apply sigmoid to get predictions
        y_pred = sigmoid(z)

        # Compute the loss
        loss = compute_loss(y, y_pred)
        loss_list.append(loss)
        if epoch % 100 == 0:
            loss_list2.append(loss)
            index.append(epoch)
            print ("Cost after iteration %i: %f" %(epoch, loss))
        # Gradients calculation
        dw = np.dot(X.T, (y_pred - y)) / m # Gradient w.r.t. weights
        db = np.sum(y_pred - y) / m        # Gradient w.r.t. bias

        # Update weights and bias
        weights -= learning_rate * dw
        bias -= learning_rate * db

    # Print Loss every 100 epochs

```

```

        if epoch % 100 == 0:
            print(f"Epoch {epoch}, Loss: {loss:.4f}")

    plt.plot(index, loss_list2)
    plt.xticks(index, rotation='vertical')
    plt.xlabel("Number of Iterarion")
    plt.ylabel("Cost")
    plt.show()

    return weights, bias, loss_list

```

```

In [59]: ##### model prediction
def predict(X, weights, bias, threshold=0.5):
    probabilities = sigmoid(np.dot(X, weights) + bias)
    return probabilities, (probabilities >= threshold).astype(int)

```

```

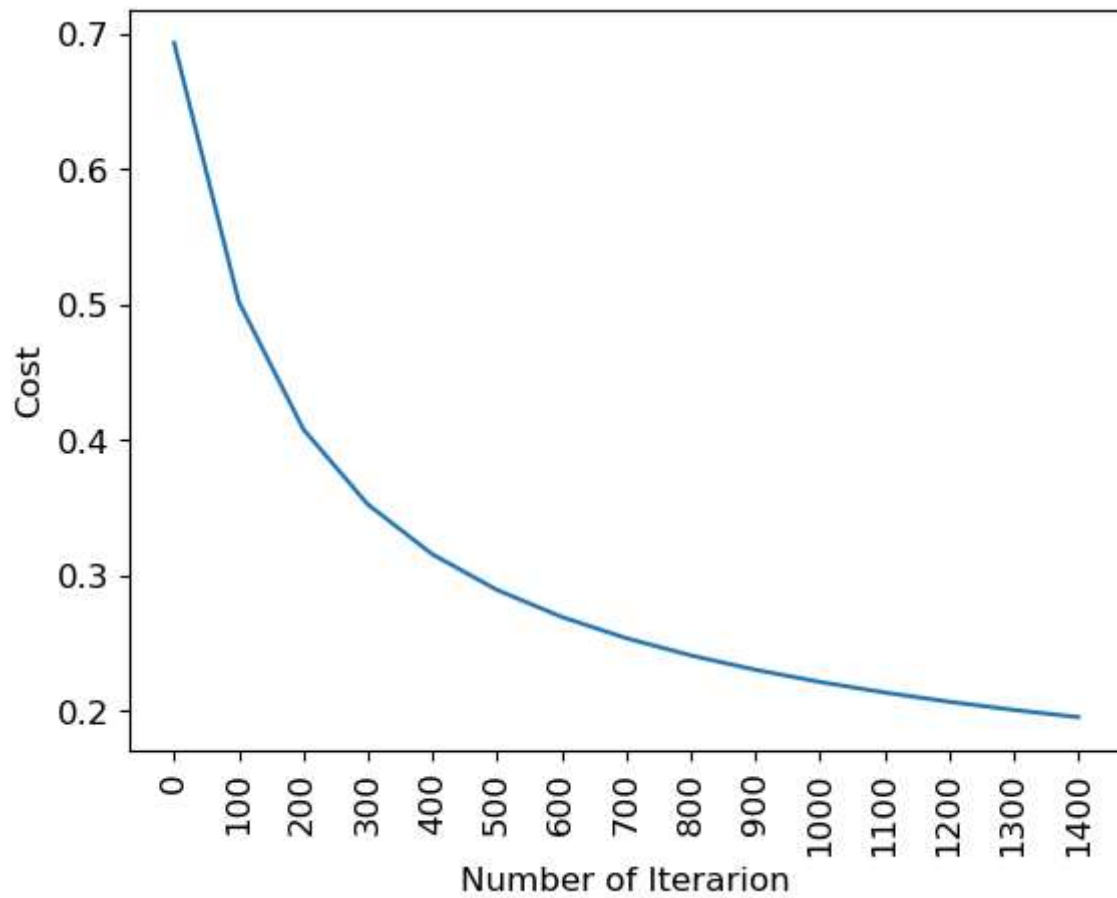
In [93]: weights, bias, loss = logistic_regression(x_train, y_train, learning_rate=0.1, epoc

```

```

Cost after iteration 0: 0.693147
Epoch 0, Loss: 0.6931
Cost after iteration 100: 0.502415
Epoch 100, Loss: 0.5024
Cost after iteration 200: 0.408058
Epoch 200, Loss: 0.4081
Cost after iteration 300: 0.352446
Epoch 300, Loss: 0.3524
Cost after iteration 400: 0.315645
Epoch 400, Loss: 0.3156
Cost after iteration 500: 0.289314
Epoch 500, Loss: 0.2893
Cost after iteration 600: 0.269406
Epoch 600, Loss: 0.2694
Cost after iteration 700: 0.253726
Epoch 700, Loss: 0.2537
Cost after iteration 800: 0.240986
Epoch 800, Loss: 0.2410
Cost after iteration 900: 0.230376
Epoch 900, Loss: 0.2304
Cost after iteration 1000: 0.221364
Epoch 1000, Loss: 0.2214
Cost after iteration 1100: 0.213583
Epoch 1100, Loss: 0.2136
Cost after iteration 1200: 0.206775
Epoch 1200, Loss: 0.2068
Cost after iteration 1300: 0.200749
Epoch 1300, Loss: 0.2007
Cost after iteration 1400: 0.195362
Epoch 1400, Loss: 0.1954

```



```
In [94]: y_pred_proba,y_pred=predict(x_test,weights,bias)
```

```
In [95]: y_test,y_pred,y_pred_proba
```

```
Out[95]: (array([0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0,
                1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0,
                0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0,
                1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1],
          dtype=int64),
         array([0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0,
                1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0,
                0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0,
                1, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1]),
         array([0.22089955, 0.94302252, 0.77537316, 0.08669849, 0.04877984,
                0.99895021, 0.99728603, 0.72816116, 0.51242851, 0.1337358 ,
                0.13849229, 0.69581088, 0.15348226, 0.53184003, 0.13333444,
                0.91464059, 0.14687007, 0.01525046, 0.00404958, 0.98347057,
                0.40481132, 0.09767545, 0.99699439, 0.02371154, 0.05701281,
                0.07926288, 0.1866041 , 0.07780257, 0.09783395, 0.96707677,
                0.0623139 , 0.06179841, 0.02182269, 0.14569493, 0.02878225,
                0.05823498, 0.46429857, 0.06085171, 0.91885101, 0.20456271,
                0.02744386, 0.68581178, 0.14077995, 0.08375229, 0.08659077,
                0.12328601, 0.01918214, 0.01694625, 0.13857154, 0.19073914,
                0.88334535, 0.9842543 , 0.33266925, 0.14375922, 0.03843487,
                0.21201424, 0.05262118, 0.99968556, 0.68404074, 0.06216004,
                0.13239926, 0.98362705, 0.99138522, 0.18948306, 0.04789653,
                0.28297061, 0.93188638, 0.99531674, 0.04370612, 0.21278713,
                0.56025266, 0.8390829 , 0.12958817, 0.88379436, 0.00891054,
                0.23470442, 0.14118216, 0.44351548, 0.03649441, 0.14869229,
                0.70965004, 0.02809926, 0.4161922 , 0.99317502, 0.68257577,
                0.88181064]))
```


```
In [96]: from sklearn.metrics import confusion_matrix
```

```
cm = confusion_matrix(y_test, y_pred)
```

```
In [97]: cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0'],
                                index=['Predict Positive:1', 'Predict Negative:0'])

sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')
```

```
Out[97]: <Axes: >
```




```
In [98]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.93	0.98	0.95	54
1	0.97	0.88	0.92	32
accuracy			0.94	86
macro avg	0.95	0.93	0.94	86
weighted avg	0.94	0.94	0.94	86


```
In [99]: plt.hist(y_pred_proba, bins = 10)
# set the title of predicted probabilities
plt.title('Histogram of predicted probabilities')
# set the x-axis limit
plt.xlim(0,1)
# set the title
plt.xlabel('Predicted probabilities')
plt.ylabel('Frequency')
```

```
Out[99]: Text(0, 0.5, 'Frequency')
```





```
In [100... from sklearn.metrics import roc_curve
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
plt.figure(figsize=(6,4))
plt.plot(fpr, tpr, linewidth=2)
plt.plot([0,1], [0,1], 'k--' )
plt.title('ROC curve for classifier')
plt.xlabel('False Positive Rate (1 - Specificity)')
plt.ylabel('True Positive Rate (Sensitivity)')
plt.show()
print(fpr)
print(tpr)
```



```
[0.      0.      0.      0.01851852 0.01851852 1.      ]
[0.      0.03125 0.875   0.875   1.      1.      ]
```

```
In [101... from sklearn.metrics import roc_auc_score

ROC_AUC = roc_auc_score(y_test, y_pred)

print('ROC AUC : {:.4f}'.format(ROC_AUC))
```

```
ROC AUC : 0.9282
```